

JQA Logistic Regression

June 18, 2022

[176]: *#Importing the Libraries*

```
import pandas as pd
import numpy as np
from sklearn import preprocessing
import matplotlib.pyplot as plt
plt.rc("font", size=14)
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import seaborn as sns
sns.set(style="white")
sns.set(style="whitegrid", color_codes=True)
```

[177]: *#Importing the Dataset*

```
df = pd.read_excel('/Users/AmrinSinghDhillon/Desktop/Book2.xlsx')
```

[178]: df

```
[178]:      is_in_play  pitch  three_plus  batter_stance  pitcher_throws  strikes  \
0              1    SL           0.0                R                L         0.0
1              1    CU           0.0                R                R         1.0
2              1    FA           0.0                R                L         2.0
3              1    SI           0.0                R                L         1.0
4              1    SL           0.0                L                R         2.0
...          ...    ...           ...                ...                ...
99995          1    CH           0.0                R                R         2.0
99996          1    CH           0.0                L                R         1.0
99997          1    FA           0.0                R                L         0.0
99998          0    FA           NaN                R                L         0.0
99999          1    FA           0.0                L                R         1.0

      balls  outs  pitch_plate_location_x  pitch_plate_location_z  \
0         1.0    2          -0.356                1.840
1         1.0    2           0.095                2.140
2         2.0    0          -0.556                2.515
3         1.0    1          -0.212                3.107
4         1.0    1          -0.248                3.437
```

```

...      ...      ...
99995      2.0      1      1.028      2.376
99996      0.0      2      -0.784      2.604
99997      0.0      1      -0.213      3.022
99998      0.0      1      -0.018      2.732
99999      2.0      2      -0.008      2.659

      pitch_initial_speed pitch_arc_break_x pitch_arc_break_z \
0      87.9      0.404316      -6.27682
1      81.0      4.821380      -13.18070
2      91.8      2.652270      -3.52130
3      92.2      5.926690      -3.88490
4      83.1      -1.288710      -9.83063
...      ...      ...
99995      74.0      -0.175697      -7.11359
99996      83.9      1.872990      -6.53433
99997      92.7      2.753500      -4.84395
99998      87.5      1.426750      -3.95007
99999      91.4      1.835650      -4.32523

      pitch_spin_rate inning pitch_per_atbat home_team_runs \
0      1025.700      3      2      1
1      2256.700      7      3      2
2      2457.170      5      5      0
3      3510.320      6      3      0
4      478.459      8      7      4
...      ...      ...
99995      1811.010      7      7      5
99996      1417.850      4      2      3
99997      1907.690      9      1      6
99998      2183.500      1      1      0
99999      1919.020      4      4      1

      away_team_runs
0      0
1      3
2      1
3      0
4      2
...      ...
99995      4
99996      4
99997      2
99998      0
99999      10

```

[100000 rows x 18 columns]

```
[179]: #Viewing the Columns
```

```
df = df.dropna()
print(df.shape)
print(list(df.columns))
```

```
(99254, 18)
```

```
['is_in_play', 'pitch', 'three_plus', 'batter_stance', 'pitcher_throws',
 'strikes', 'balls', 'outs', 'pitch_plate_location_x', 'pitch_plate_location_z',
 'pitch_initial_speed', 'pitch_arc_break_x', 'pitch_arc_break_z',
 'pitch_spin_rate', 'inning', 'pitch_per_atbat', 'home_team_runs',
 'away_team_runs']
```

```
[184]: #Turning Categorical Variables into Binary
```

```
cat_vars=['pitch','three_plus','batter_stance','pitcher_throws']
for var in cat_vars:
    cat_list='var'+ '_' +var
    cat_list = pd.get_dummies(df[var], prefix=var)
```

```
[189]: df=df1
```

```
[190]: df
```

```
[190]:
```

	is_in_play	pitch	three_plus	batter_stance	pitcher_throws	strikes	\
0	1	SL	0.0	R	L	0.0	
1	1	CU	0.0	R	R	1.0	
2	1	FA	0.0	R	L	2.0	
3	1	SI	0.0	R	L	1.0	
4	1	SL	0.0	L	R	2.0	
...	
99994	1	SI	0.0	L	L	2.0	
99995	1	CH	0.0	R	R	2.0	
99996	1	CH	0.0	L	R	1.0	
99997	1	FA	0.0	R	L	0.0	
99999	1	FA	0.0	L	R	1.0	

	balls	outs	pitch_plate_location_x	pitch_plate_location_z	...	\
0	1.0	2	-0.356	1.840	...	
1	1.0	2	0.095	2.140	...	
2	2.0	0	-0.556	2.515	...	
3	1.0	1	-0.212	3.107	...	
4	1.0	1	-0.248	3.437	...	
...	
99994	3.0	0	-0.014	3.006	...	
99995	2.0	1	1.028	2.376	...	
99996	0.0	2	-0.784	2.604	...	

99997	0.0	1		-0.213		3.022	...
99999	2.0	2		-0.008		2.659	...

	pitch_FA	pitch_FC	pitch_SI	pitch_SL	three_plus_0.0	three_plus_1.0	\
0	0	0	0	1	1	0	
1	0	0	0	0	1	0	
2	1	0	0	0	1	0	
3	0	0	1	0	1	0	
4	0	0	0	1	1	0	
...	
99994	0	0	1	0	1	0	
99995	0	0	0	0	1	0	
99996	0	0	0	0	1	0	
99997	1	0	0	0	1	0	
99999	1	0	0	0	1	0	

	batter_stance_L	batter_stance_R	pitcher_throws_L	pitcher_throws_R
0	0	1	1	0
1	0	1	0	1
2	0	1	1	0
3	0	1	1	0
4	1	0	0	1
...
99994	1	0	1	0
99995	0	1	0	1
99996	1	0	0	1
99997	0	1	1	0
99999	1	0	0	1

[99254 rows x 30 columns]

```
[191]: cat_vars=['three_plus','batter_stance','pitcher_throws','pitch']
df_vars=df.columns.values.tolist()
to_keep=[i for i in df_vars if i not in cat_vars]
```

```
[192]: #Viewing New Categories with Binary Variables

df=df[to_keep]
df.columns.values
```

```
[192]: array(['is_in_play', 'strikes', 'balls', 'outs', 'pitch_plate_location_x',
        'pitch_plate_location_z', 'pitch_initial_speed',
        'pitch_arc_break_x', 'pitch_arc_break_z', 'pitch_spin_rate',
        'inning', 'pitch_per_atbat', 'home_team_runs', 'away_team_runs',
        'pitch_CH', 'pitch_CU', 'pitch_FA', 'pitch_FC', 'pitch_SI',
        'pitch_SL', 'three_plus_0.0', 'three_plus_1.0', 'batter_stance_L',
        'batter_stance_R', 'pitcher_throws_L', 'pitcher_throws_R'],
      dtype=object)
```

```
dtype=object)
```

```
[193]: #Identifying the Dependent and Independent Variables
```

```
X = df.loc[:, df.columns != 'is_in_play']  
y = df.loc[:, df.columns == 'is_in_play']
```

```
[195]: #Splitting the Data into Train and Test Sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
↳random_state=0)  
logreg = LogisticRegression()  
logreg.fit(X_train, y_train)
```

```
/Users/amrinsinghdhillon/opt/anaconda3/lib/python3.8/site-  
packages/sklearn/utils/validation.py:1111: DataConversionWarning: A column-  
vector y was passed when a 1d array was expected. Please change the shape of y  
to (n_samples, ), for example using ravel().  
y = column_or_1d(y, warn=True)  
/Users/amrinsinghdhillon/opt/anaconda3/lib/python3.8/site-  
packages/sklearn/linear_model/_logistic.py:444: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[195]: LogisticRegression()
```

```
[196]: #Printing Accuracy of the Classifier
```

```
y_pred = logreg.predict(X_test)  
print('Accuracy of logistic regression classifier on test set: {:.2f}'.  
↳format(logreg.score(X_test, y_test)))
```

Accuracy of logistic regression classifier on test set: 0.68

```
[197]: #Viewing the Classification Report
```


```
from sklearn.metrics import classification_report  
print(classification_report(y_test, y_pred))
```

```
precision    recall  f1-score   support
```

0	0.69	0.27	0.39	11254
1	0.68	0.93	0.78	18523
accuracy			0.68	29777
macro avg	0.68	0.60	0.58	29777
weighted avg	0.68	0.68	0.63	29777

[198]: *#Creating the ROC Curve*

```
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
logit_roc_auc = roc_auc_score(y_test, logreg.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, logreg.predict_proba(X_test)[:,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



[]: